

Beyond Model-Free Reinforcement Learning

Solving Optimal Decision Making Problems from a Different Perspective

Matthia Sabatelli

December 17, 2021

Recap

Last week we have seen ...

Model-Free Reinforcement Learning

- How to construct algorithms when the state $|\mathcal{S}|$ and action spaces $|\mathcal{A}|$ are large
- Function Approximators
- Linear and Non-Linear Functions
- Deep Reinforcement Learning
- Experience Replay $D = \langle s_t, a_t, r_t, s_{t+1} \rangle$

Today's Agenda

- ① Policy Gradient Methods
- ② Beyond Model-Free Learning
- ③ Beyond Online Learning
- ④ Beyond Markov Decision Processes

Policy Gradient Methods

Throughout this course we have constantly seen how important **value functions** are:

- They correspond to the knowledge of the agent
- They help us understanding the environment
- But most importantly they define the policy π that the agent follows

⇒ The most **powerful** value function is the state-action value function $Q(s, a)$ since:

$$\pi^*(s) = \arg \max_{a \in \mathcal{A}} Q^*(s, a) \text{ for all } s \in \mathcal{S}$$

Policy Gradient Methods

Our **goal** so far has always been that of learning $Q(s, a)$:

1. We can do this in a tabular fashion
2. Or we can generalize this process with a function approximator
3. We can do this in an *on-policy* or *off-policy* fashion

⇒ **First** we learn a value function, and **then** we derive the policy π

⇒ We have **never** seen how to learn π directly!

Policy Gradient Methods

Action value based methods such as Q-Learning, SARSA and $QV(\lambda)$ are very powerful algorithms but have some important **limitations**:

1. Are restricted to discrete action spaces
2. Work alongside carefully designed exploration-exploitation strategies
3. Are unable to learn an optimal policy which is stochastic
4. Sometimes learning $\pi(a|s; \theta)$ is easier than learning $Q(s, a; \theta)$

Policy Gradient Methods

There is a family of techniques which tries to learn $\pi(a|s; \theta)$ directly:

Policy Gradient Methods

⇒ They learn a **parametrized policy** that learns how to select actions without having to consult a value function:

$$\pi(a|s; \theta) = \Pr \{a_t = a, s_t = s ; \theta_t = \theta\}$$

The parameters θ usually correspond to the weights of a neural network

Policy Gradient Methods

Training policy gradient methods significantly differs from training a Deep-Q Network:

⇒ The idea of last week's methods was that of **minimizing** a certain quantity called the TD-error:

$$\mathcal{L}(\theta) = \mathbb{E}_{\langle \cdot \rangle \sim U(D)} \left[\left(r_t + \gamma \max_{a \in \mathcal{A}} Q(s_{t+1}, a; \theta^-) - Q(s_t, a_t; \theta) \right)^2 \right]$$

⇒ Policy Gradient methods instead seek to **maximize** some scalar performance measure $J(\theta)$ and update the parameters via gradient ascent!

$$\theta_{t+1} = \theta + \alpha \widehat{\nabla J(\theta_t)}$$

Policy Gradient Methods

It is also possible to learn $\pi(a|s; \theta)$ in combination with a value function!

⇒ These algorithms come with the name of **Actor-Critic** methods:

- It can be hard to learn a policy π directly
- We would like to tell our agent learning who is learning π how good its policy is
- To do so we can use the state-value function $V^\pi(s)$

Policy Gradient Methods

How do Actor-Critic methods intuitively work?

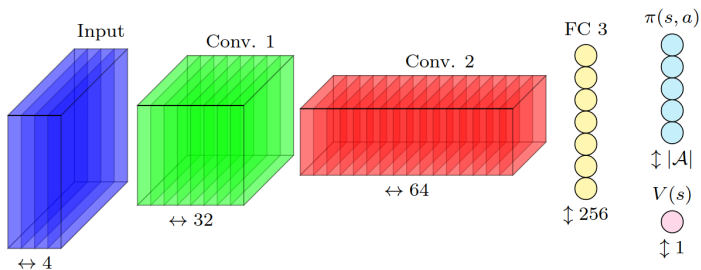


Figure: Image courtesy of Van de Wolfshaar (2017)

Policy Gradient Methods

How do we train this network?

$$\begin{aligned}\theta_{t+1} &= \theta + \alpha \left(r_t + \gamma V(s_{t+1}; \phi) - V(s_t; \phi) \right) \frac{\nabla \pi(a_t | s_t; \theta_t)}{\pi(a_t | s_t; \theta_t)} \\ &= \theta_t + \alpha \delta_t \frac{\nabla \pi(a_t | s_t; \theta_t)}{\pi(a_t | s_t; \theta_t)}\end{aligned}$$

⇒ Note that a value function (ϕ) **helps** learning the policy parameters θ but is **not required** for action selection purposes.

Policy Gradient Methods

⇒ Nowadays actor-critic algorithms have become **very popular** in Deep Reinforcement Learning:

1. They are **theoretically motivated** thanks to the *policy gradient theorem* (see page 325 of the book)
2. The critic can technically learn **any value function**
3. Can be massively parallelized (see A3C algorithm)

⇒ However, compared to action-value based methods, actor-critic algorithms are **less well understood** (maybe because learning a policy π is still more complex than learning a value function?)

Beyond Model-Free Learning

⇒ Actor-Critic algorithms, just like action value based methods are also **model-free** Reinforcement Learning algorithms. As a result we have never even attempted learning the transition function \mathcal{P} of the Markov Decision Process \mathcal{M} .

- Recall that the \mathcal{P} and \mathcal{R} components of \mathcal{M} are usually called the **model** of the environment
- If they are known we can use Dynamic Programming algorithms like **value iteration** :)
- However, in the typical Reinforcement Learning scenario this is **never** the case :(

Beyond Model-Free Learning

What to do?

- In Model-Based Reinforcement Learning the goal is to **learn the model** of the environment through experience!
- The **idea** is to learn a function that comes in the following form: $f(s_t, a_t) = s_{t+1}$
- If learned f would give us $p(s_{t+1}|s_t, a_t)$
- We can do something similar for learning $p(r_t|a_t, s_t)$

⇒ The task of learning a model of the environment corresponds to a **supervised learning** problem!

Beyond Model-Free Learning

⇒ The overall learning strategy is very **simple**:

1. We start with a random policy $\pi(a_t|s_t)$
2. This policy results in a dataset of trajectories
 $\mathcal{D} = \{(s, a, s')_i\}$
3. We learn the dynamics of the model by minimizing
 $\sum_i \|f(s_i, a_i) - s'_i\|^2$
4. Plan through the model and go back to step 2.

Beyond Model-Free Learning

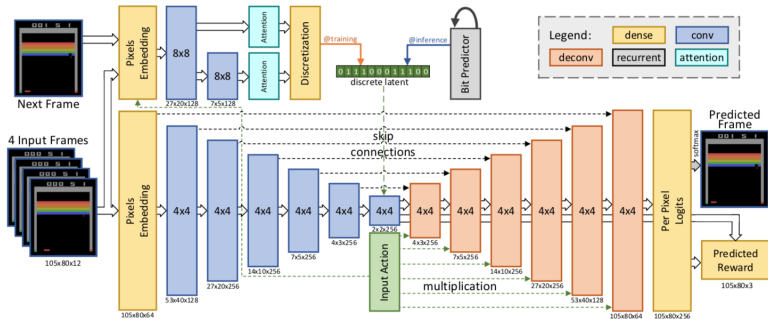


Figure: Image courtesy of Kaiser et al. (2020)

Beyond Model-Free Learning




Figure: Image courtesy of Ha et al. (2019)

Beyond Model-Free Learning

Pros & Cons of Model-Based Reinforcement Learning:

- Some argue that learning $p(s_{t+1}|s_t, a_t)$ is easier than learning $Q^\pi(s, a)$ or $\pi(a|s)$ ✓
- If some dynamics of the environment are known it is easy to include them in the learning process ✓
- As it is essentially supervised learning it scales to a wide range of other applications (GANs and VAEs) ✓
- Is computationally very expensive ✗
- What if the model is wrong? ✗

Beyond Online Learning

Throughout our lectures we have always assumed that the agent-environment interaction:

1. Can last as long as we want (possibly **infinite**)
2. Is **inexpensive**
3. Moving from s_t to s_{t+1} is **fast**
4. We have as many trajectories $\tau\langle s_t, a_t, r_t, s_{t+1} \rangle$ as we want at our disposal for learning

⇒ However, for many **practical applications** this is not the case!

Beyond Online Learning

Imagine the following situations:

- Clinical trials
- Self-driving cars
- Drones
- Finance

⇒ These are all examples for which it is **very costly** to interact with the environment!

We do have access to **some data** of the environment but not to the environment itself.

Beyond Online Learning

Sometimes we have to train Reinforcement Learning algorithms with **limited** data samples


Batch Reinforcement Learning

⇒ We have a small dataset of trajectories $\tau\langle s_t, a_t, r_t, s_{t+1}\rangle$, collected by an unknown policy π and wish to learn an **optimal value function**.

The learning agent is **not allowed** to gather any new experience!

Beyond Online Learning

Online Reinforcement Learning



Offline / Batch RL




Figure: Image courtesy of Swazinna et al. (2021)

Beyond Online Learning

⇒ How do we **train** an agent in the Batch Reinforcement Learning context?

1. We start from a **dataset** D of trajectories τ collected by policy π'
2. We train our favourite model-free algorithm on D until **convergence**: $Q^{\pi'}(s, a)$
3. We can **deploy** the learned state-action value function to the environment and hope for the best

Where is the **tricky** part?

- Our dataset of trajectories is very small **✗**
- Do you trust the Q function that you have learned on D ? **✗**
- We can only collect D a few times **✗**

Beyond Online Learning

⇒ Batch Reinforcement Learning poses some very interesting **challenges** to the community:

1. It forces us to think in **practical terms**: we do not always have access to a simulation of the environment
2. It questions the **reliability** of the mathematical tools we are using
3. We need to reason about **uncertainty** even more: what if D is only representative of a very small portion of \mathcal{M}
4. Opens the door to **transfer-learning** training strategies

Beyond Markov Decision Processes

We have always modeled the **environment** as a Markov Decision Process:

$$\mathcal{M}\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, r \rangle$$

and assumed that the entire state space \mathcal{S} is **visible** to the agent.



Beyond Markov Decision Processes

⇒ However there are many practical applications where it is **impossible** for an agent to observe some states of the environment



How do we **formalize** such an environment?

Beyond Markov Decision Processes

⇒ We **generalize** the concept of a MDP to a setting where once we have taken action a_t we either can, or **cannot** observe s_{t+1} : Partially Observable Markov Decision Process (POMDP)

$$\mathcal{M}\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, r, \Omega, \mathcal{O}, \gamma \rangle$$

We add two new components to the MDP:

1. Ω is a set of observations
2. \mathcal{O} are the conditional observation probabilities

Ω gives us a **hint** about the state the agent is visiting, but not the true state itself!

Beyond Markov Decision Processes

When dealing with a POMDP an agent deals with the **uncertainty** of the state space \mathcal{S} by maintaining a **belief** b over the visited states:

- b represents a probability distribution $\mathcal{P}(\mathcal{S})$
- $b(s)$ denotes the probability $P(\mathcal{S} = s)$ under the current belief state
- Throughout interaction with the environment b is updated

\Rightarrow The **goal** becomes to learn $\pi^*(b)$

The End !

In Memoriam: Marco A. Wiering (1971-2021)

